Machine Vision

October 19, 2018

1 Lecture 1

Bernoulli Distribution

$$Pr(x) = \lambda^x(1 - \lambda)^{1-x}, \lambda \in= [0, 1], x \in \{0, 1\}$$

$$Pr(x) = Bern_x[\lambda]$$

Beta Distribution

$$Pr(\lambda) = \frac{\Gamma[\alpha + \beta]}{\Gamma[\alpha]\Gamma[\beta]}\lambda^{\alpha-1}(1 - \lambda)^{\beta-1}, \alpha, \beta > 0$$

$$\Gamma(z) = \int_0^\infty t^{z-1}c^{-t}dt = (z - 1)!$$

$$E[\lambda] = \frac{\alpha}{\alpha + \beta}$$

$$B(p, q) = \frac{q - 1}{p + q + 1}B(p, q - 1)$$

α, β decide the coin fact λ

Categorical Distribution

$$Pr(x = k) = \lambda_k$$

$$Pr(x) = Cat_x[\lambda]$$

Dirichlet Distribution

$$Pr(\lambda_1 \ldots \lambda_K) = \frac{\Gamma[\sum_{k=1}^{K}]\alpha_k}{\prod_{k=1}^{K}\Gamma[\alpha_k]}\prod_{k=1}^{K}\lambda_k^{\alpha_k-1}$$

$$Pr(\lambda_1 \ldots \lambda_K) = \text{Dir}_{\lambda_1\ldots\lambda_K}[\alpha_1, \alpha_2 \ldots, \alpha_K]$$

Univariate Normal Distribution

$$Pr(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp[-0.5(x-\mu)^2/\sigma^2]$$

$$Pr(x) = \text{Norm}_x[\mu, \sigma^2]$$

Normal Inverse Gamma Distribution

$$Pr(\mu, \sigma^2) = \frac{\sqrt{\gamma}\beta^\alpha}{\sigma\sqrt{2\pi}\Gamma[\alpha]} \left(\frac{1}{\sigma^2}\right)^{\alpha+1} \exp[-\frac{2\beta + \gamma(\delta-\mu)^2}{2\sigma^2}]$$

$$Pr(\mu, \sigma^2) = \text{NormInvGam}_{\mu,\sigma^2}[\alpha, \beta, \gamma, \delta]$$

Multivariate Normal Distribution

$$Pr(\mathbf{x}) = \frac{1}{(2\pi)^{D/2}|\sum|^{1/2}} \exp[-0.5(\mathbf{x}-\mu)^T \sum{}^{-}1(\mathbf{x}-\mu)]$$

Normal Inverse Wishart

$$Pr(\mu, \Sigma) = \frac{\gamma^{D/2}|\Psi|^{\alpha/2}|\Sigma|^{-\frac{\alpha+D+2}{2}}}{(2\pi)^{D/2}2^{\frac{\alpha D}{2}}\Gamma_D(\frac{\alpha}{2})} \exp\{-\frac{1}{2}(Tr(\Psi\Sigma^{-1})) + \gamma(\mu-\delta)^T\Sigma^{-1}(\mu-\delta)\}$$

Conjugate Distribution and Conjugate prior

Conjugate Distribution is between prior and posterior

Prior is the conjugate prior of the likelihood function.

2 Fitting model

maximum likelihood

Fitting

$$\begin{aligned}
\hat{\theta} &= argmax_{(\theta)}[Pr(\mathbf{x_{1...I}}|\theta)] \\
&= argmax_{(\theta)}[\prod_{i=1}^{I} Pr(\mathbf{x_i}|\theta)]
\end{aligned}$$

Predictive Density

Evaluate new data point $\mathbf{x}^*$ under probability distribution $Pr(\mathbf{x}^*|\hat{\theta})$ with best parameter.

maximum a posteriori

Fitting

$$
\begin{aligned}
\hat{\theta} &= argmax_{(\theta)}[Pr(\theta|\mathbf{x_{1...I}})] \\
&= argmax_{(\theta)}\left[\frac{Pr(\mathbf{x_{1...I}}|\theta)Pr(\theta)}{Pr(\mathbf{x_{1...I}})}\right] \\
&= argmax_{(\theta)}\left[\frac{\prod_{i=1}^{I} Pr(\mathbf{x_i}|\theta)Pr(\theta)}{Pr(\mathbf{x_{1...I}})}\right] \\
\hat{\theta} &= argmax_{(\theta)}\left[Pr(\mathbf{x_i}|\theta)Pr(\theta)\right]
\end{aligned}
$$

Predictive

Evaluate new data point $\mathbf{x}^*$ under probability distribution $Pr(\mathbf{x}^*|\hat{\theta})$ with best parameter.

bayesian approach

Fitting

$$
Pr(\theta|\mathbf{x_{1...I}}) = \frac{(\prod_{i=1}^{I} Pr(\mathbf{x_i}|\theta))Pr(\theta)}{Pr(\mathbf{x_{1...I}})}
$$

The difference between bayesian approach and MAP is that MAP takes the maximum value, while bayesian approach takes the distribution.

Predictive

$$
Pr(\mathbf{x}^*|\mathbf{x_{1...I}}) = \int Pr(\mathbf{x}^*|\theta)Pr(\theta|\mathbf{x_{1...I}})d\theta
$$

Confusion: the formula should be $\int Pr(\mathbf{x}^*|\theta, \mathbf{x_{1...I}})Pr(\theta|\mathbf{x_{1...I}})d\theta$. Given the θ, it considers $\mathbf{x_{1...I}}$ and x^* are independent.

Multivariate Normal Distribution

If $\mathbf{x_1}, \mathbf{x_2} \ldots \mathbf{x_n}$ are independent, the covariance matrix would be

$$
\Sigma = \begin{bmatrix}
\sigma_1^2 & 0 & \ldots & 0 \\
0 & \sigma_2^2 & \ldots & 0 \\
\vdots & \vdots & \ddots & 0 \\
0 & 0 & \ldots & \sigma_n^2
\end{bmatrix}
$$

Therefore, while $\mathbf{x_1}, \mathbf{x_2} \ldots \mathbf{x_n}$ are dependent, the covariance matrix could be decomposed into rotation matrix and diagonal:

$$\Sigma_{full} = \mathbf{R}^T \Sigma'_{diag} \mathbf{R}$$

Marginal Distribution

$$u_i = u_i$$

$$\Sigma_i = \Sigma_{ii}$$

Conditional Distribution

$$u_{i|j} = u_i + \Sigma_{ij}\Sigma_{jj}^{-1}(x_j - u_j)$$

$$\Sigma_{i|j} = \Sigma_{jj} - \Sigma_{ij}^T\Sigma_{ii}^{-1}\Sigma_{ij}$$

Product of two normals

$$\begin{aligned}
\text{Norm}_{\mathbf{x}}[\mathbf{a}, \mathbf{A}]\text{Norm}_{\mathbf{x}}[\mathbf{b}, \mathbf{B}] &= k \cdot \text{Norm}_{\mathbf{x}}[(\mathbf{A}^{-1} + \mathbf{B}^{-1})^{-1}(\mathbf{A}^{-1}\mathbf{a} + \mathbf{B}^{-1}\mathbf{b}), (\mathbf{A}^{-1} + \mathbf{B}^{-1})^{-1}] \\
k &= \text{Norm}_{\mathbf{a}}[\mathbf{b}, \mathbf{A} + \mathbf{B}]
\end{aligned}$$

change of variables

$$\text{Norm}_{\mathbf{x}}[\mathbf{Ay} + \mathbf{b}, \Sigma] = k \cdot \text{Norm}_{\mathbf{y}}[\mathbf{A'x} + \mathbf{b'}, \Sigma']$$

where

$$\begin{aligned}
\mathbf{A'} &= \Sigma' A^T \Sigma^{-1} \\
b' &= -\Sigma' A^T \Sigma^{-1} b \\
\Sigma &= (A^T \Sigma^{-1} A)^{-1}
\end{aligned}$$

Learning and Inference

The observe measured data, $\mathbf{x}$

Draw inference from it about the state of world, $\mathbf{w}$

If $\mathbf{w}$ is continuous, call this regression.

If $\mathbf{w}$ is discrete, call this classification.

To compute the probability distribution $Pr(\mathbf{w}|\mathbf{x})$, we need: a model(relates visual data $\mathbf{x}$ and $\mathbf{w}$, the relationships depends on parameter θ), a learning algorithm(fits parameter θ from paired training examples $\mathbf{x_i}$, $\mathbf{w_i}$), an inference algorithm (use model to return $Pr(\mathbf{w}|\mathbf{x})$ given new observed data $\mathbf{x}$)

Types of Model

1. Model contingency of the world on the data $Pr(\mathbf{w}|\mathbf{x})$ (Discriminative models)

1. Choose an appropriate from form for $Pr(\mathbf{w})$

2. Make parameters a function of $\mathbf{x}$

3. Function takes parameters θ that define its shape.

Inference: evaluate $\mathbf{Pr(w|x)}$

2. Model joint occurrence of the world and data $Pr(\mathbf{x}, \mathbf{w})$ Generative models

1. COncatenate $\mathbf{x}$ and $\mathbf{w}$ to make $\mathbf{z} = [\mathbf{x^T w^T}]$

2. Model of pdf of $\mathbf{z}$

3. Pdf takes parameter θ that define its shape

Inference: compute $Pr(\mathbf{w}|\mathbf{x})$ using Bayes rule.

$$Pr(\mathbf{w}|\mathbf{x}) = \frac{Pr(\mathbf{x}, \mathbf{w})}{Pr(\mathbf{x})} = \frac{Pr(\mathbf{x}, \mathbf{w})}{\int Pr(\mathbf{x}, \mathbf{w})d\mathbf{w}}$$

3. Model contingency of data on the world $Pr(\mathbf{x}|\mathbf{w})$ (Generative models)

1. Choose an appropriate form for $Pr(\mathbf{x})$

2. Make parameters a function of $\mathbf{w}$

3. Function takes parameter θ that define its shape.

Inference: define prior $Pr(\mathbf{w})$ and then compute $Pr(\mathbf{w}|\mathbf{x})$ using Bayes' rule.

$$Pr(\mathbf{w}|\mathbf{x}) = \frac{Pr(\mathbf{x}|\mathbf{w})Pr(\mathbf{w})}{\int Pr(\mathbf{x}|\mathbf{w})\mathbf{Pr(w)dw}}$$

Bessel correction

$s^2 = (\frac{n}{n-1})s_n^2$ working this later.

Learning and inference

Mixture of Gaussian

$$Pr(\mathbf{x}|\theta) = \sum_{k=1}^{K} \lambda_k \text{Norm}_{\mathbf{x}}[\mu_k, \Sigma_k]$$

Hidden variables

$$
\begin{aligned}
Pr(\mathbf{x}) &= \int Pr(\mathbf{x}, \mathbf{h})d\mathbf{h} \\
Pr(\mathbf{x}|\theta) &= \int Pr(\mathbf{x}, \mathbf{h}|\theta)d\mathbf{h} \\
\hat{\theta} &= \text{argmax}_\theta \left[\sum_{i=1}^{\mathbf{I}} log[\int Pr(\mathbf{x}_i, \mathbf{h}_i|\theta)d\mathbf{h}_i] \right] \\
B[\{q_i(\mathbf{h}_i)\}, \theta] &= \sum_{i=1}^{\mathbf{I}} \int q_i(\mathbf{h}_i) \log[\frac{Pr(\mathbf{x}, \mathbf{h}_i|\theta)}{q_i(\mathbf{h}_i)}]d\mathbf{h}_{1...I} \leq \sum_{i=1}^{\mathbf{I}} log[\int Pr(\mathbf{x}_i, \mathbf{h}_i|\theta)d\mathbf{h}_i]
\end{aligned}
$$